

Intro To Statistical Computing

Final Project

MATH 208

Authors: *Ling Fei Zhang, Brandon Ma*

Instructor: *Russell J. Steele*
Department: *Statistics*
Date: *November 30, 2022*

Contents

1	Set up	2
2	Task 1	3
2.1	part a)	3
2.2	part b)	9
2.3	part c)	12
3	Task 2	15
3.1	Part a)	15
3.2	Part b)	20

1 Set up

```
library(readxl)
library(here)
library(matlib)
library(formatR)
library(gridExtra)
library(tidyverse)
library(tidyr)
knitr::opts_chunk$set(tidy.opts=list(width.cutoff=60),
                      tidy=TRUE,
                      echo=TRUE,
                      # comment=NA,
                      message=FALSE,
                      warning=FALSE)
cpu_gpu_data<-read_csv("chip_dataset.csv")
```

```
names(cpu_gpu_data)
```

```
## [1] "ID"           "Product"      "Type"
## [4] "Release Date" "Process Size (nm)" "TDP (W)"
## [7] "Die Size (mm^2)" "Transistors (million)" "Freq (MHz)"
## [10] "Foundry"      "Vendor"       "FP16 GFLOPS"
## [13] "FP32 GFLOPS"  "FP64 GFLOPS"
```

```
head(cpu_gpu_data)
```

```
## # A tibble: 6 x 14
##   ID Product      Type Release~1 Proce~2 TDP (~3 Die S~4 Trans~5 Freq ~6 Foundry
##   <dbl> <chr>      <chr> <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl> <chr>
## 1 0 AMD Athlo~ CPU 2007-0~ 65 45 77 122 2200 Unknown
## 2 1 AMD Athlo~ CPU 2018-0~ 14 35 192 4800 3200 Unknown
## 3 2 Intel Cor~ CPU 2020-0~ 10 28 NA NA 2600 Intel
## 4 3 Intel Xeo~ CPU 2013-0~ 22 80 160 1400 1800 Intel
## 5 4 AMD Pheno~ CPU 2011-0~ 45 125 258 758 3700 Unknown
## 6 5 Intel Xeo~ CPU 2013-0~ 22 95 160 1400 2400 Intel
## # ... with 4 more variables: Vendor <chr>, `FP16 GFLOPS` <dbl>,
## # `FP32 GFLOPS` <dbl>, `FP64 GFLOPS` <dbl>, and abbreviated variable names
## # 1: `Release Date`, 2: `Process Size (nm)`, 3: `TDP (W)`,
## # 4: `Die Size (mm^2)`, 5: `Transistors (million)`, 6: `Freq (MHz)`
```

2 Task 1

2.1 part a)

```
# Numerical summary in tables
summary <- cpu_gpu_data %>%
  select(Type, "Process Size (nm)", "TDP (W)", "Die Size (mm^2)",
         "Transistors (million)", "Freq (MHz)") %>%
  pivot_longer(cols = all_of(c("Process Size (nm)", "TDP (W)",
                                "Die Size (mm^2)", "Transistors (million)", "Freq (MHz)")),
               names_to = "Characteristics") %>%
  group_by(Type, Characteristics) %>%
  summarise_all(list(Min = min, Max = max, Med = median, IQR = IQR),
                na.rm = TRUE)

summary

## # A tibble: 10 x 6
## # Groups:   Type [2]
##   Type Characteristics      Min    Max   Med    IQR
##   <chr> <chr>          <dbl> <dbl> <dbl> <dbl>
## 1 CPU   Die Size (mm^2)         1    684   149   108
## 2 CPU   Freq (MHz)             600   4700  2400  1000
## 3 CPU   Process Size (nm)        7    180    32    76
## 4 CPU   TDP (W)                 1    400    65    60
## 5 CPU   Transistors (million)    37 19200   410 1086
## 6 GPU   Die Size (mm^2)         6    826   148   155
## 7 GPU   Freq (MHz)             100  2321   600   438
## 8 GPU   Process Size (nm)        0    250    40    62
## 9 GPU   TDP (W)                 2    900    50   91.5
## 10 GPU  Transistors (million)    8 54200   716 2590

# Graphical summary
p1 <- cpu_gpu_data %>%
  ggplot(aes(x = Type, y = `Process Size (nm)`, fill = Type)) +
  stat_boxplot(geom = "errorbar", width = 0.25) + geom_boxplot(na.rm = TRUE) +
  ylab("Process Size (nm)") + xlab("Type") + ggtitle("Box plot of Process Size (nm) by Type")
p2 <- cpu_gpu_data %>%
  ggplot(aes(x = `Process Size (nm)`, col = Type)) + geom_density(size = 1,
    na.rm = TRUE) + xlab("Process Size (nm)") + ylab("Density") +
  ggtitle("Density plot of Process Size (nm) by Type")

grid.arrange(p1, p2)

p1 <- cpu_gpu_data %>%
  ggplot(aes(x = Type, y = `TDP (W)`, fill = Type)) + stat_boxplot(geom = "errorbar",
    width = 0.25) + geom_boxplot(na.rm = TRUE) + ylab("TDP (W)") +
  xlab("Type") + ggtitle("Box plot of TDP (W) by Type")
p2 <- cpu_gpu_data %>%
  ggplot(aes(x = `TDP (W)`, col = Type)) + geom_density(size = 1,
    na.rm = TRUE) + xlab("TDP (W)") + ylab("Density") + ggtitle("Density plot of TDP (W) by Type")

grid.arrange(p1, p2)

p1 <- cpu_gpu_data %>%
  ggplot(aes(x = Type, y = `Die Size (mm^2)`, fill = Type)) +
```

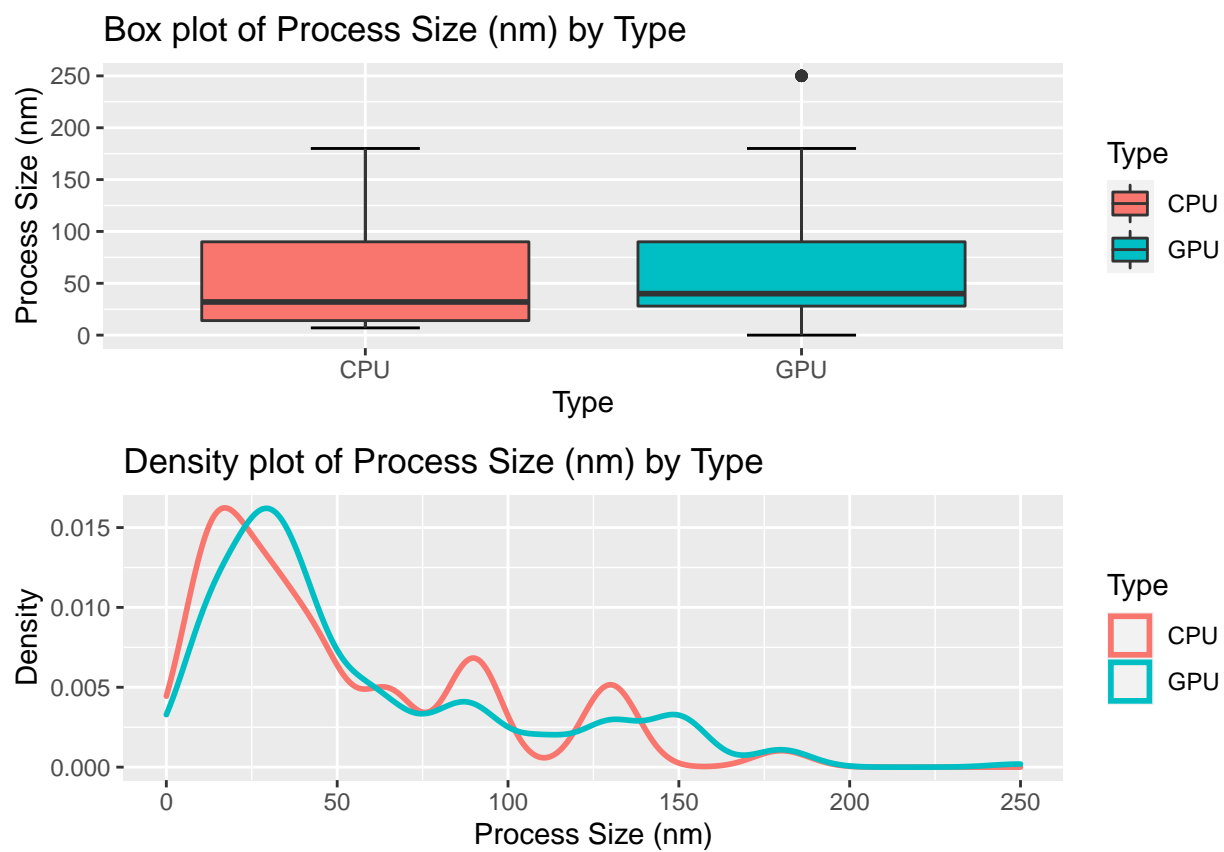


Figure 1: Process Size Data

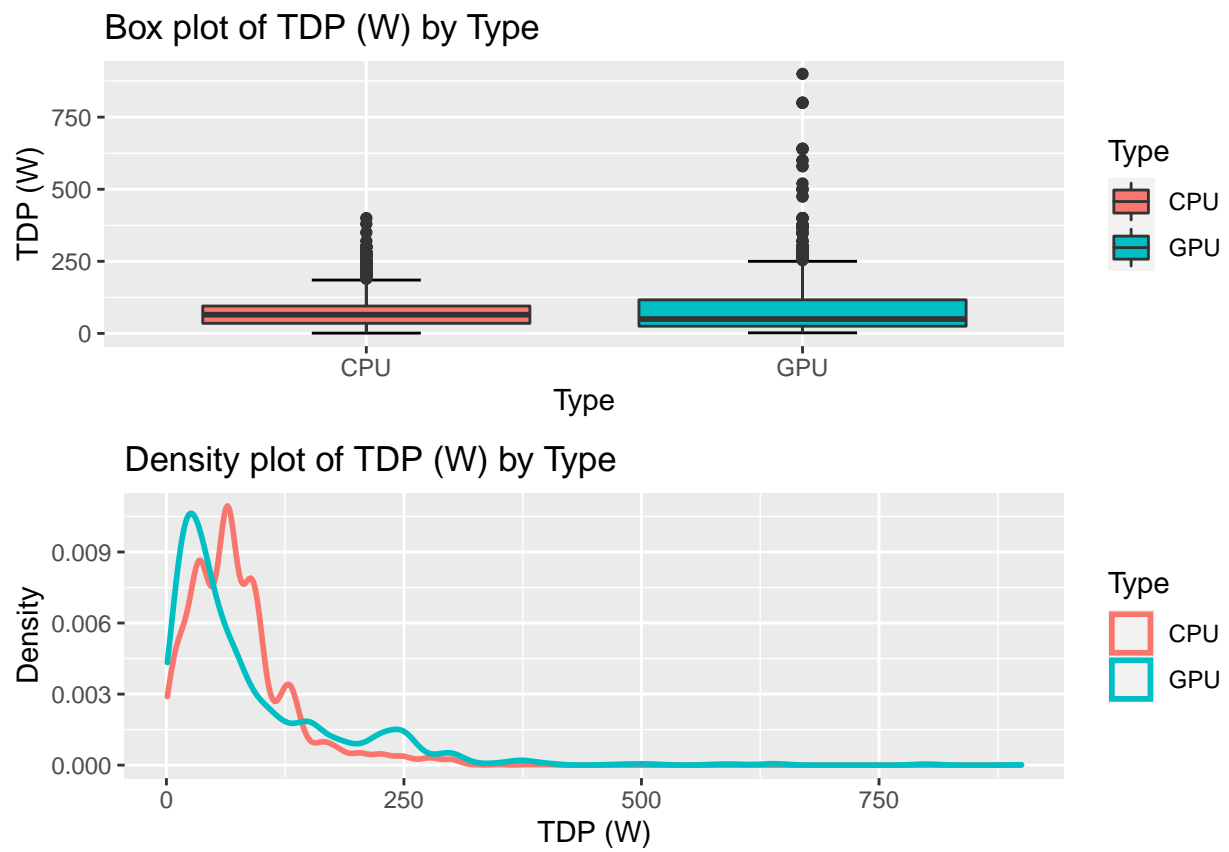


Figure 2: TDP Data

```

stat_boxplot(geom = "errorbar", width = 0.25) + geom_boxplot(na.rm = TRUE) +
  ylab("Die Size (mm^2)") + xlab("Type") + ggtitle("Box plot of Die Size by Type")
p2 <- cpu_gpu_data %>%
  ggplot(aes(x = `Die Size (mm^2)`, col = Type)) + geom_density(size = 1,
    na.rm = TRUE) + xlab("Die Size (mm^2)") + ylab("Density") +
    ggtitle("Density plot of Die Size (mm^2) by Type")

grid.arrange(p1, p2)

```

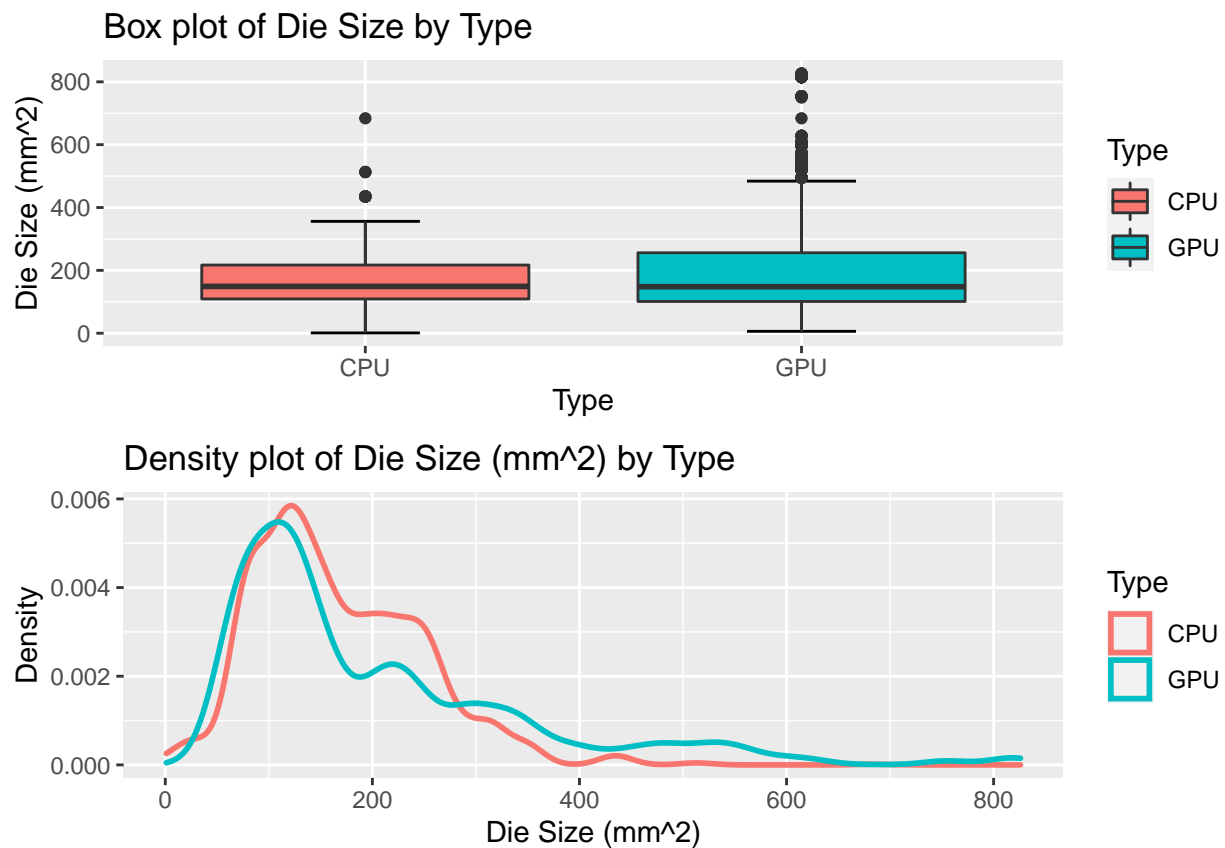


Figure 3: Die Size Data

```

p1 <- cpu_gpu_data %>%
  ggplot(aes(x = Type, y = `Transistors (million)`, fill = Type)) +
  stat_boxplot(geom = "errorbar", width = 0.25) + geom_boxplot(na.rm = TRUE) +
  ylab("Transistors (million)") + xlab("Type") + ggtitle("Box plot of Transistors (million) by Type")
p2 <- cpu_gpu_data %>%
  ggplot(aes(x = `Transistors (million)`, col = Type)) + geom_density(size = 1,
    na.rm = TRUE) + xlab("Transistors (million)") + ylab("Density") +
    ggtitle("Density plot of Transistors (million) by Type")

grid.arrange(p1, p2)

p1 <- cpu_gpu_data %>%
  ggplot(aes(x = Type, y = `Freq (MHz)`, fill = Type)) + stat_boxplot(geom = "errorbar",
    width = 0.25) + geom_boxplot(na.rm = TRUE) + ylab("Freq (MHz)") +

```

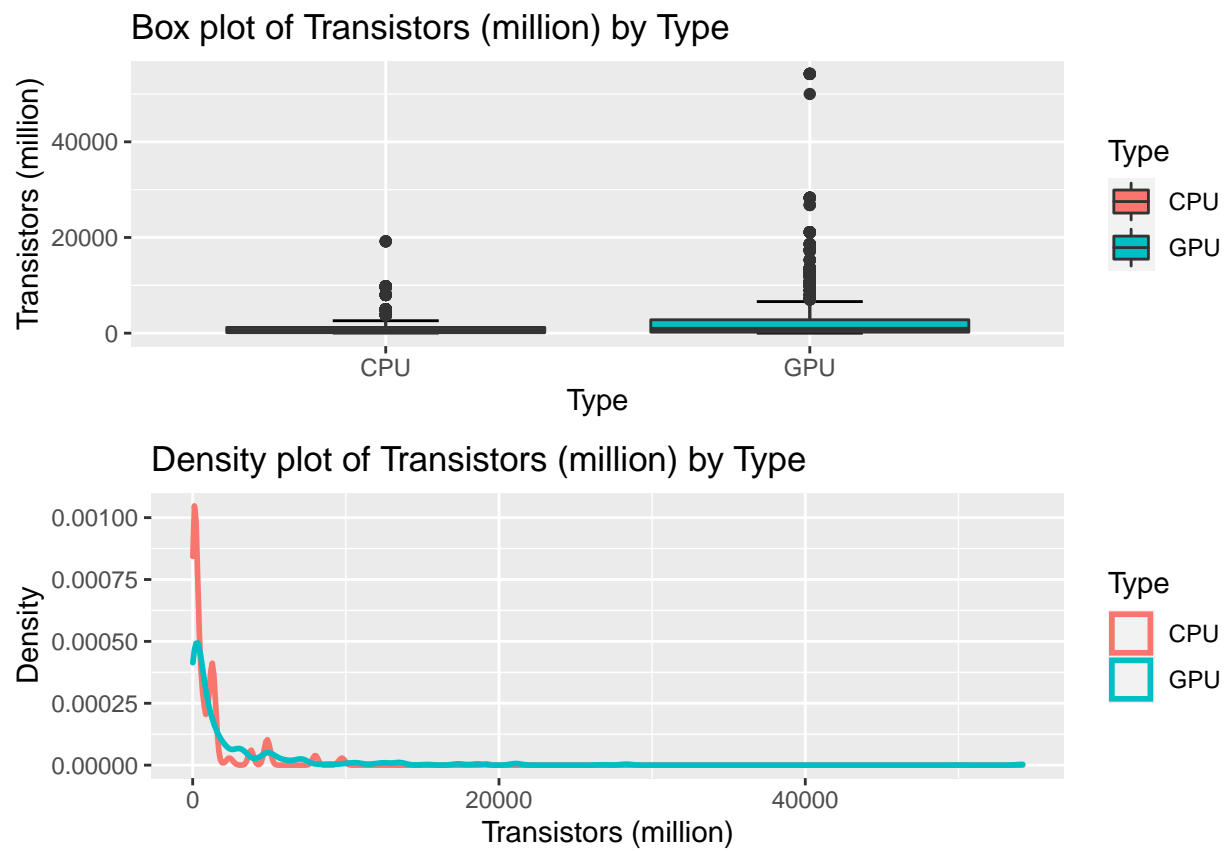


Figure 4: Transistors Data


```

xlab("Type") + ggtitle("Box plot of Freq (MHz) by Type")
p2 <- cpu_gpu_data %>%
  ggplot(aes(x = `Freq (MHz)`, col = Type)) + geom_density(size = 1,
    na.rm = TRUE) + xlab("Freq (MHz)") + ylab("Density") + ggtitle("Density plot of Freq (MHz) by Type")
grid.arrange(p1, p2)

```

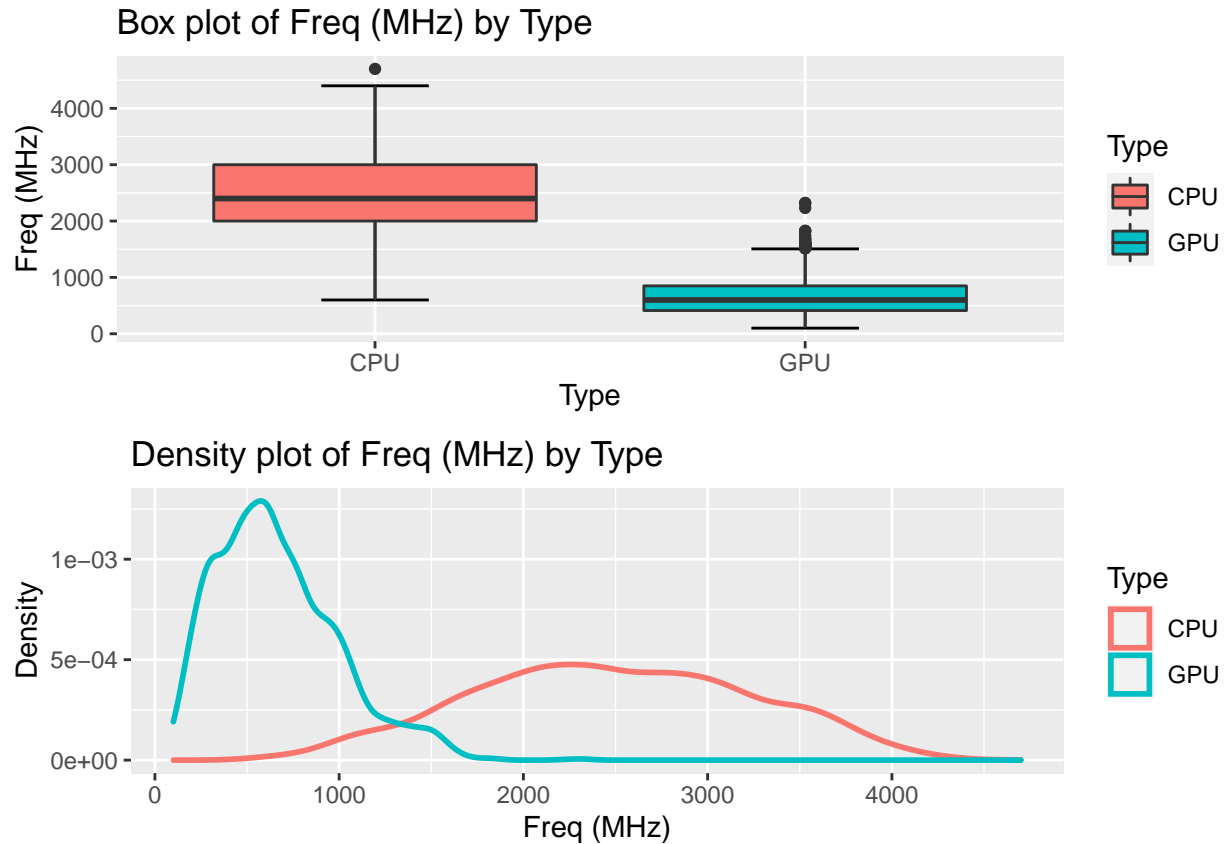


Figure 5: Frequency Data

2.1.1 Process Size

From the numerical summary, we can see that CPUs has a minimum of 7, maximum of 180, median of 32 and IQR of 76 and that GPUs has a minimum of 0, maximum of 250, median of 40 and IQR of 62. The characteristics of the distribution can also be observed from Figure 1. From the boxplot, we can see that the central location and spread of both types are very similar. However, the GPUs have an outlier at the point 250, where the max comes from. The density plot tells us that the skewness of both types are similar.

2.1.2 TDP

From the numerical summary, we can see that CPUs has a minimum of 1, maximum of 400, median of 65 and IQR of 60 and that GPUs has a minimum of 2, maximum of 900, median of 50 and IQR of 91.5. The characteristics of the distribution can also be observed from Figure 2. From the boxplot, we can see that the central locations are very similar, whereas the spread of the distribution is quite different as GPUs reach up to 900. The density plot tells us that the skewness of the two types are quite similar.

2.1.3 Die Size

From the numerical summary, we can see that CPUs has a minimum of 1, maximum of 684, median of 149 and IQR of 108 and that GPUs has a minimum of 6, maximum of 826, median of 148 and IQR of 155. The characteristics of the distribution can also be observed from Figure 3. From the boxplot, we can see that the central locations are very similar, with a slightly greater spread with the GPUs. The density plot tells us that the skewness of the two types are very similar.

2.1.4 Transistors

From the numerical summary, we can see that CPUs has a minimum of 37, maximum of 19200, median of 410 and IQR of 1086 and that GPUs has a minimum of 8, maximum of 54200, median of 716 and IQR of 2590. The characteristics of the distribution can also be observed from Figure 4. From the boxplot, we can see that the central locations are somewhat similar. However, the spread of the GPUs is much greater than the spread of the CPUs. This can be easily observed, as the maximum number of transistors in a CPU is 19200 in contrast to 54200 in GPU. We can also observe a few outliers in the GPUs from the boxplot. From the density plot, we can see that the skewness of the distributions are also quite different. We can see that proportional to the size of the distribution, CPUs tend to have a smaller number of transistors compared to GPUs, as the skewness is much higher when the number of transistors is small.

2.1.5 Frequency

From the numerical summary, we can see that CPUs has a minimum of 600, maximum of 4700, median of 2400 and IQR of 1000 and that GPUs has a minimum of 100, maximum of 2321, median of 600 and IQR of 438. The characteristics of the distribution can also be observed from Figure 5. From the boxplot, we can see that the central locations are very different. This is clear as the medians are 2400 and 600 for CPUs and GPUs respectively. The density plot tells us that both the skewness and the spread of the two types are very different. CPUs have a bigger spread, but a smaller skewness. On the other hand, GPUs have a smaller spread, but a greater skewness.

2.2 part b)

```
# Numerical Summary
table(cpu_gpu_data["Foundry"])

## Foundry
##      GF      IBM    Intel    NEC Renesas Samsung    Sony    TSMC    UMC Unknown
##    265      3    1390      2      1      60      10    2178     79     866

vendor_vs_foundry <- cpu_gpu_data %>%
  group_by(Vendor, Type) %>%
  summarise(`unique foundry` = n_distinct(Foundry))
vendor_vs_foundry

## # A tibble: 7 x 3
## # Groups:   Vendor [5]
##   Vendor Type `unique foundry`
##   <chr> <chr>         <int>
## 1 AMD    CPU             3
## 2 AMD    GPU             4
## 3 ATI    GPU             5
## 4 Intel  CPU             1
## 5 Intel  GPU             2
## 6 NVIDIA GPU             5
## 7 Other  GPU             4
```

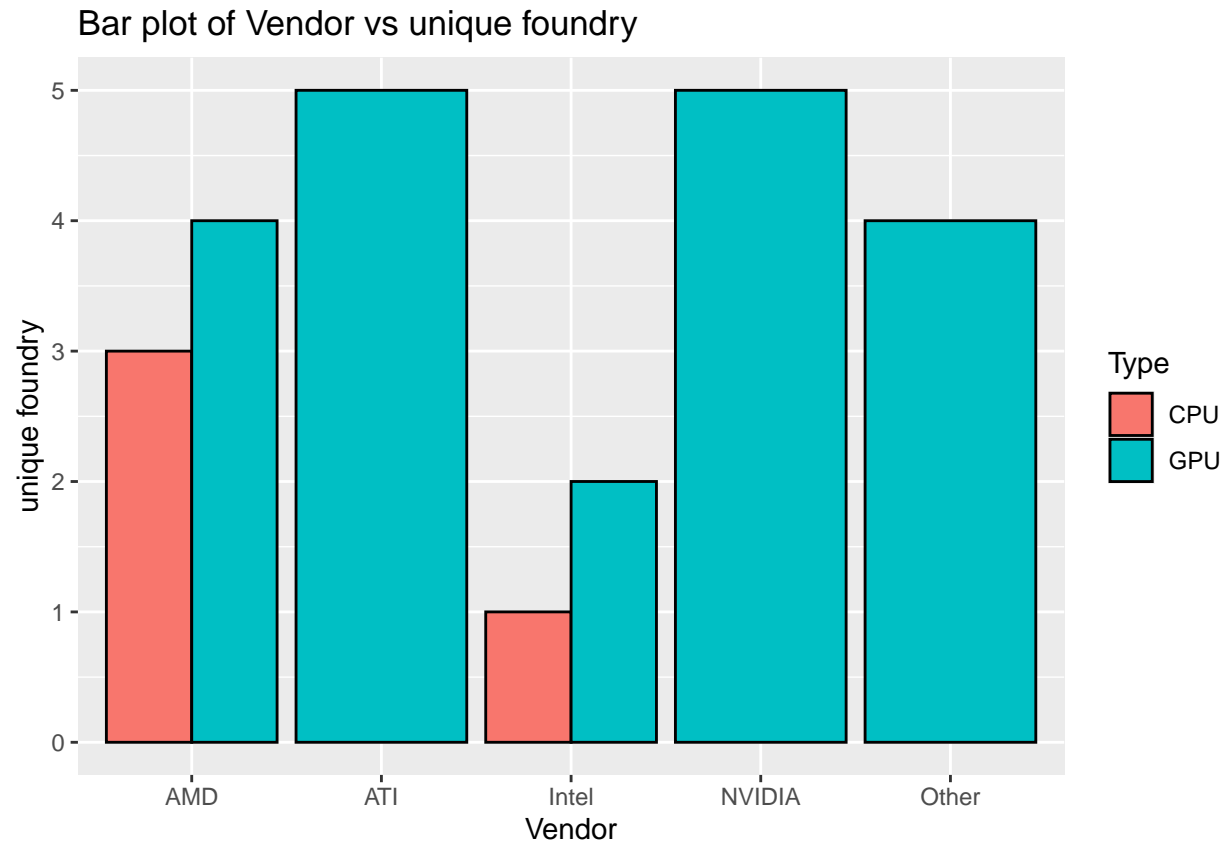
We can see that the only vendor that has a unique foundry is Intel, but that is only specific to the CPU processors. For GPUs, even Intel has 2 distinct foundries. All other vendors have more than 1 distinct foundry, regardless of the type of the processor.

```
foundry_vs_vendor <- cpu_gpu_data %>%
  group_by(Foundry, Type) %>%
  summarise(`unique vendor` = n_distinct(Vendor))
foundry_vs_vendor
```


```
## # A tibble: 14 x 3
## # Groups:   Foundry [10]
##   Foundry Type `unique vendor`
##   <chr>   <chr>         <int>
## 1 GF      CPU             1
## 2 GF      GPU             1
## 3 IBM     GPU             1
## 4 Intel   CPU             1
## 5 Intel   GPU             1
## 6 NEC     GPU             1
## 7 Renesas GPU             1
## 8 Samsung GPU             2
## 9 Sony    GPU             2
## 10 TSMC    CPU             1
## 11 TSMC    GPU             4
## 12 UMC     GPU             3
## 13 Unknown CPU             1
## 14 Unknown GPU             4
```

Here, we can see from the second table that many of the foundries have a unique vendor. We should keep in mind that IMB, NEC, Renesas, Samsung, Sony and UMC are all foundries that have a small count. We observe that in general, each foundry like to stay with a small number of unique vendors. Additionally, we observe that **all CPU foundries have a unique vendor**. This information can also be observed from the graphs below, perhaps even better.

```
# Graphical Summary
vendor_vs_foundry %>%
  ggplot(aes(x = Vendor, y = `unique foundry`, group = Type,
    fill = Type)) + geom_bar(col = "black", stat = "identity",
    position = "dodge") + ggtitle("Bar plot of Vendor vs unique foundry")
```




```
foundry_vs_vendor %>%  
  ggplot(aes(x = Foundry, y = `unique vendor`, group = Type,  
            fill = Type)) + geom_bar(col = "black", stat = "identity",  
            position = "dodge") + ggtitle("Bar plot of Foundry vs unique vendor")
```




2.3 part c)

```
# Graph
cpu_gpu_data %>%
  ggplot(aes(x = `Die Size (mm^2)`, y = `TDP (W)`, col = Type)) +
  geom_point() + xlab("Die Size (mm^2)") + ylab("TDP (W)") +
  ggtitle("Die Size vs TDP") + geom_smooth(method = "lm")
```



We can also observe the graphs separately, as follows:

```
cpu_gpu_data %>%  
  ggplot(aes(x = `Die Size (mm^2)`, y = `TDP (W)`, col = Type)) +  
  geom_point() + facet_wrap(~Type) + xlab("Die Size (mm^2)") +  
  ylab("TDP (W)") + ggtitle("Die Size vs TDP") + geom_smooth(method = "lm",  
    col = "black")
```



In both graphs, we can see that the association between Die Size and TDP does not really depend on Type. It's worth noting that the two features in GPUs have a greater spread, whereas the association in CPUs are much more concentrated. However, we can observe that both types of processors follow a similar regression line. We can also observe the association numerically.

```
# Numerical summary in tables
summary[c(1, 4, 6, 9), ]
```

```
## # A tibble: 4 x 6
## # Groups:   Type [2]
##   Type Characteristics    Min    Max    Med   IQR
##   <chr> <chr>          <dbl> <dbl> <dbl> <dbl>
## 1 CPU   Die Size (mm^2)      1    684   149  108
## 2 CPU   TDP (W)              1    400    65   60
## 3 GPU   Die Size (mm^2)      6   826   148  155
## 4 GPU   TDP (W)              2   900    50  91.5
```

From the numerical table above, we can also observe that the medians of both characteristics are similar, regardless of the processor type (CPU or GPU).

3 Task 2

3.1 Part a)

3.1.1 CPU

```
foundry_vs_year <- cpu_gpu_data[cpu_gpu_data$Type == "CPU", ] %>%  
  group_by(Foundry, year = format(as.Date(na.omit(`Release Date`),  
    format = "%Y-%m-%d"), "%Y")) %>%  
  summarise(total_count = n())
```


```
foundry_vs_year <- na.omit(foundry_vs_year)
```

```
foundry_vs_year
```

```
## # A tibble: 49 x 3  
## # Groups:   Foundry [4]  
##   Foundry year total_count  
##   <chr>   <chr>      <int>  
## 1 GF      2017          39  
## 2 GF      2018          29  
## 3 GF      2019          18  
## 4 GF      2020           7  
## 5 Intel   2000           3  
## 6 Intel   2001          35  
## 7 Intel   2002          26  
## 8 Intel   2003          32  
## 9 Intel   2004          46  
## 10 Intel  2005          59  
## # ... with 39 more rows
```

```
foundry_vs_year %>%  
  ggplot(aes(x = year, y = total_count, group = Foundry, fill = Foundry,  
    width = 0.7)) + geom_bar(col = "black", stat = "identity",  
    position = "dodge", na.rm = TRUE) + ggtitle("Total number of processors  
relased by year") +  
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5,  
    hjust = 1))
```


Total number of processors
relased by year



```
vendor_vs_year <- cpu_gpu_data[cpu_gpu_data$Type == "CPU", ] %>%
  group_by(Vendor, year = format(as.Date(na.omit(`Release Date`),
    format = "%Y-%m-%d"), "%Y")) %>%
  summarise(total_count = n())
```


```
vendor_vs_year <- na.omit(vendor_vs_year)
```

```
vendor_vs_year
```

```
## # A tibble: 44 x 3
## # Groups:   Vendor [2]
##   Vendor year total_count
##   <chr> <chr>      <int>
## 1 AMD    2000         14
## 2 AMD    2001         51
## 3 AMD    2002         13
## 4 AMD    2003         41
## 5 AMD    2004         90
## 6 AMD    2005         95
## 7 AMD    2006         66
## 8 AMD    2007         32
## 9 AMD    2008         67
## 10 AMD   2009         52
## # ... with 34 more rows
```

```
vendor_vs_year %>%
  ggplot(aes(x = year, y = total_count, group = Vendor, fill = Vendor,
```

```
width = 0.7)) + geom_bar(col = "black", stat = "identity",
  position = "dodge", na.rm = TRUE) + ggtitle("Total number of processors
  released by year") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
    hjust = 1))
```



3.1.2 GPU

```
foundry_vs_year <- cpu_gpu_data[cpu_gpu_data$Type == "GPU", ] %>%
  group_by(Foundry, year = format(as.Date(na.omit(`Release Date`),
    format = "%Y-%m-%d"), "%Y")) %>%
  summarise(total_count = n())
```

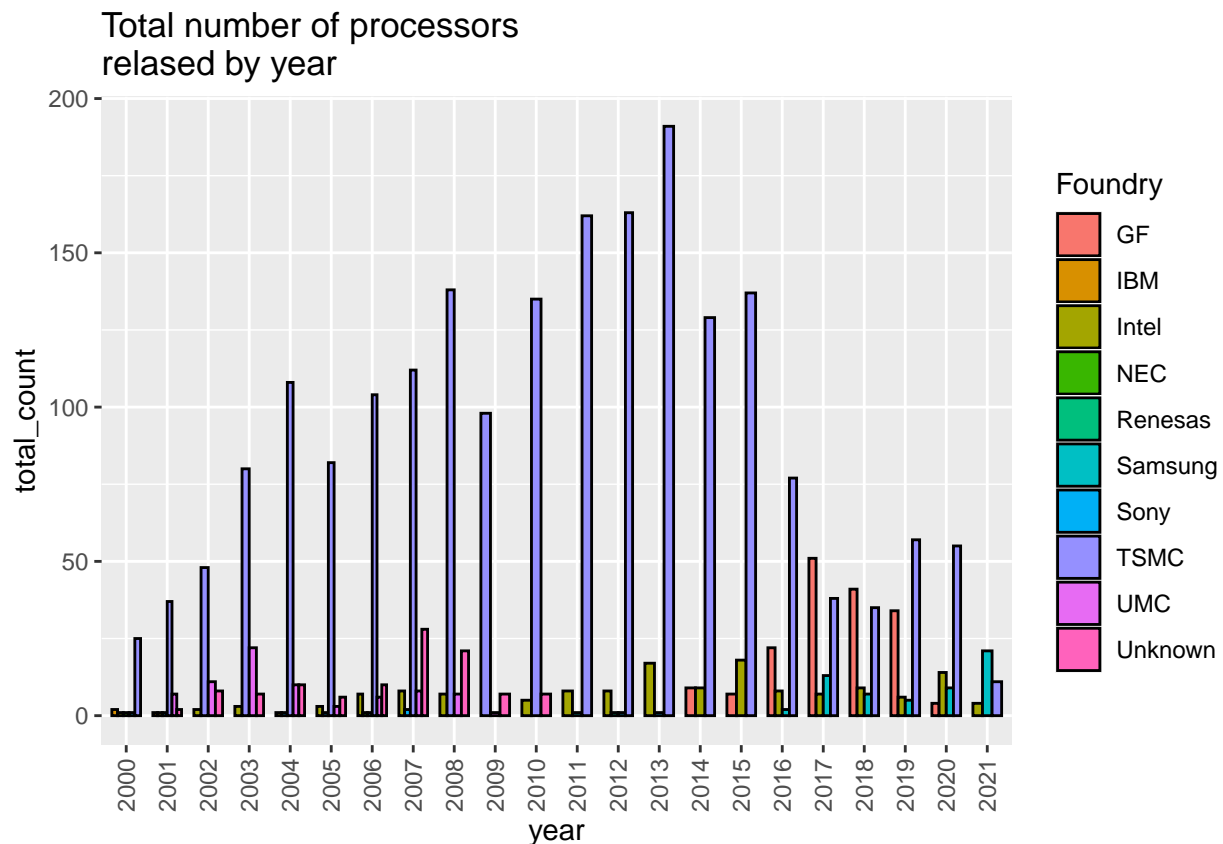
```
foundry_vs_year <- na.omit(foundry_vs_year)
```

```
foundry_vs_year
```

```
## # A tibble: 89 x 3
## # Groups:   Foundry [10]
##   Foundry year total_count
##   <chr>   <chr>      <int>
## 1 GF     2014         9
## 2 GF     2015         7
## 3 GF     2016        22
## 4 GF     2017        51
## 5 GF     2018        41
```

```
## 6 GF      2019      34
## 7 GF      2020       4
## 8 IBM     2000       2
## 9 IBM     2001       1
## 10 Intel  2000       1
## # ... with 79 more rows
```

```
foundry_vs_year %>%
  ggplot(aes(x = year, y = total_count, group = Foundry, fill = Foundry,
    width = 0.7)) + geom_bar(col = "black", stat = "identity",
    position = "dodge", na.rm = TRUE) + ggtitle("Total number of processors
    relased by year") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
    hjust = 1))
```



```
vendor_vs_year <- cpu_gpu_data[cpu_gpu_data$Type == "GPU", ] %>%
  group_by(Vendor, year = format(as.Date(na.omit(`Release Date`),
    format = "%Y-%m-%d"), "%Y")) %>%
  summarise(total_count = n())
```

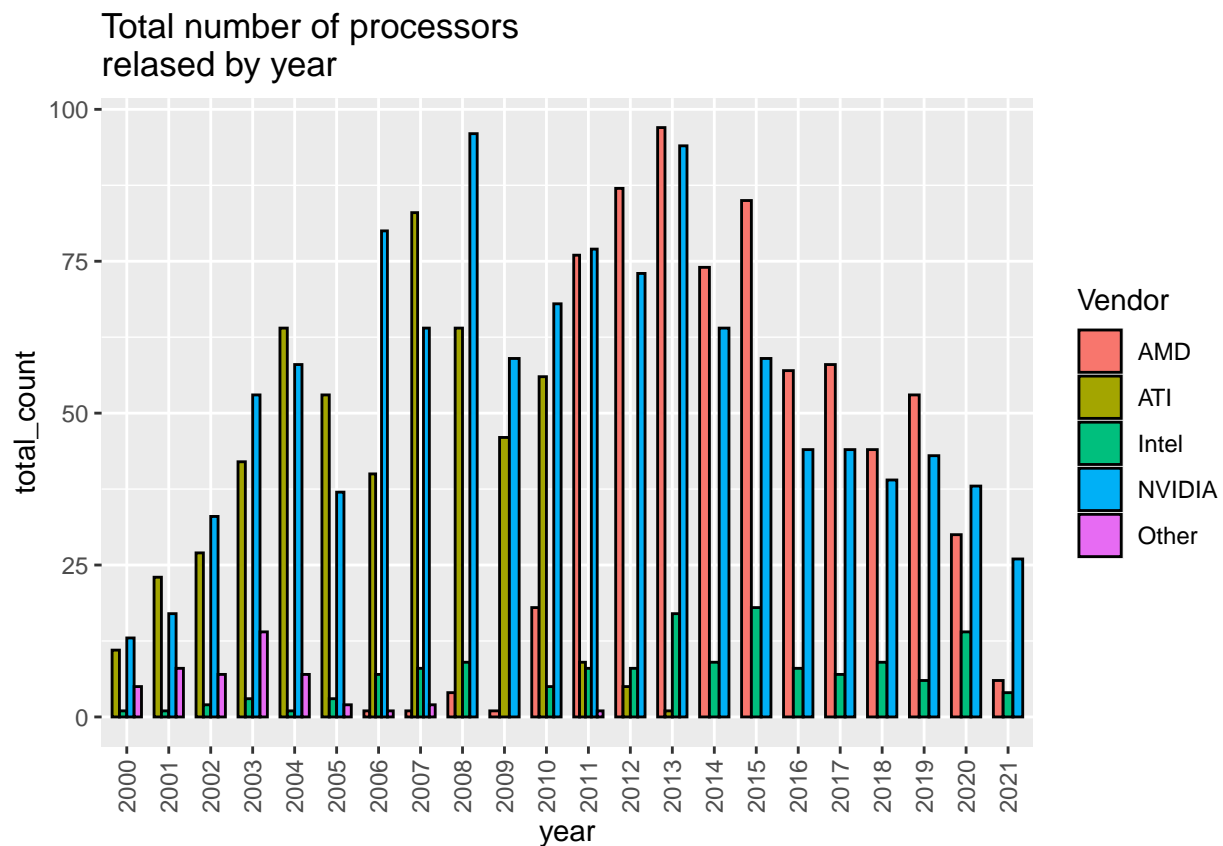
```
vendor_vs_year <- na.omit(vendor_vs_year)
```

```
vendor_vs_year
```

```
## # A tibble: 82 x 3
## # Groups:   Vendor [5]
##   Vendor year total_count
```

```
##      <chr>  <chr>      <int>
## 1 AMD      2006         1
## 2 AMD      2007         1
## 3 AMD      2008         4
## 4 AMD      2009         1
## 5 AMD      2010        18
## 6 AMD      2011        76
## 7 AMD      2012        87
## 8 AMD      2013        97
## 9 AMD      2014        74
## 10 AMD     2015        85
## # ... with 72 more rows
```

```
vendor_vs_year %>%
  ggplot(aes(x = year, y = total_count, group = Vendor, fill = Vendor,
    width = 0.7)) + geom_bar(col = "black", stat = "identity",
    position = "dodge", na.rm = TRUE) + ggtitle("Total number of processors
    relased by year") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
    hjust = 1))
```



Over the years, we have seen a general increase in release of processors from 2000 to around 2012, then a general decrease from 2012 to 2021. For CPUs, Intel was the leading foundry and vendor for most of the years. For GPUs, TSMC was the top foundry. ATI was the leading GPU vendor 2010, then AMD and NVIDIA took over.

3.2 Part b)

3.2.1 Expected Number of transistors per microchip

```
cpu_data <- cpu_gpu_data[cpu_gpu_data$Type == "CPU", ] %>%
  group_by(`Transistors (million)`, year = format(as.Date(na.omit(`Release Date`),
    format = "%Y-%m-%d"), "%Y"))
curr <- mean(cpu_data[cpu_data$year == 2000, ]$`Transistors (million)`,
  na.rm = TRUE)
CPU_exp_nb_trans <- c(curr)
years <- seq(2000, 2021, by = 1)
for (i in 1:21) {
  curr <- curr * sqrt(2)
  CPU_exp_nb_trans <- append(CPU_exp_nb_trans, curr)
}
```

3.2.1.1 CPU

```
gpu_data <- cpu_gpu_data[cpu_gpu_data$Type == "GPU", ] %>%
  group_by(`Transistors (million)`, year = format(as.Date(na.omit(`Release Date`),
    format = "%Y-%m-%d"), "%Y"))
curr <- mean(gpu_data[gpu_data$year == 2000, ]$`Transistors (million)`,
  na.rm = TRUE)
GPU_exp_nb_trans <- c(curr)
for (i in 1:21) {
  curr <- curr * sqrt(2)
  GPU_exp_nb_trans <- append(GPU_exp_nb_trans, curr)
}
```

3.2.1.2 GPU

3.2.2 Observed number of transistors per microchip

```
cpu_data <- cpu_gpu_data[cpu_gpu_data$Type == "CPU", ] %>%
  group_by(`Transistors (million)`, year = format(as.Date(na.omit(`Release Date`),
    format = "%Y-%m-%d"), "%Y"))
CPU_obs_nb_trans <- c()

for (i in 0:21) {
  curr <- mean(cpu_data[cpu_data$year == 2000 + i, ]$`Transistors (million)`,
    na.rm = TRUE)
  CPU_obs_nb_trans <- append(CPU_obs_nb_trans, curr)
}
```

3.2.2.1 CPU

```
gpu_data <- cpu_gpu_data[cpu_gpu_data$Type == "GPU", ] %>%
  group_by(`Transistors (million)`, year = format(as.Date(na.omit(`Release Date`),
    format = "%Y-%m-%d"), "%Y"))
GPU_obs_nb_trans <- c()

for (i in 0:21) {
```

```

curr <- mean(gpu_data[gpu_data$year == 2000 + i, ]$`Transistors (million)`,
             na.rm = TRUE)
GPU_obs_nb_trans <- append(GPU_obs_nb_trans, curr)
}

```

3.2.2.2 GPU

```

result <- data.frame(years, CPU_exp_nb_trans, CPU_obs_nb_trans,
                     GPU_exp_nb_trans, GPU_obs_nb_trans)
result

```


3.2.2.3 All together

##	years	CPU_exp_nb_trans	CPU_obs_nb_trans	GPU_exp_nb_trans	GPU_obs_nb_trans
## 1	2000	38.00000	38.00000	25.03704	25.03704
## 2	2001	53.74012	65.19767	35.40772	42.53191
## 3	2002	76.00000	51.84615	50.07407	49.77612
## 4	2003	107.48023	89.31507	70.81543	73.44444
## 5	2004	152.00000	104.21324	100.14815	124.62609
## 6	2005	214.96046	137.46575	141.63087	164.60227
## 7	2006	304.00000	269.11864	200.29630	215.09244
## 8	2007	429.92092	252.11111	283.26174	295.70470
## 9	2008	608.00000	335.61290	400.59259	509.34337
## 10	2009	859.84185	533.86087	566.52348	613.83019
## 11	2010	1216.00000	530.23776	801.18519	959.77931
## 12	2011	1719.68369	919.41333	1133.04695	1224.07059
## 13	2012	2432.00000	1450.51852	1602.37037	1685.98266
## 14	2013	3439.36738	1522.49587	2266.09391	2165.05612
## 15	2014	4864.00000	1591.91667	3204.74074	2298.16541
## 16	2015	6878.73477	1382.41667	4532.18782	2727.20588
## 17	2016	9728.00000	5269.85714	6409.48148	4482.54839
## 18	2017	13757.46953	5170.67857	9064.37564	5622.32323
## 19	2018	19456.00000	6626.77273	12818.96296	8429.37500
## 20	2019	27514.93907	4252.48077	18128.75128	7262.18750
## 21	2020	38912.00000	6597.50000	25637.92593	14899.23077
## 22	2021	55029.87814	9800.00000	36257.50256	19382.00000

```

ggplot(data = result, aes(years)) + geom_point(aes(y = CPU_exp_nb_trans,
color = "CPU Expected")) + geom_point(aes(y = CPU_obs_nb_trans,
color = "CPU Observed")) + geom_point(aes(y = GPU_exp_nb_trans,
color = "GPU Expected")) + geom_point(aes(y = GPU_obs_nb_trans,
color = "GPU Observed"))

```



3.2.2.4 Graphically

Hence, from the table and plot above, we can say that Moore's Law is no longer valid.